

Comparison of AWS Database Services

| Service | Type | Cost (Estimated) | Performance | When to Use | Key Differences |
|------------|---------------------------|--|--|---|--|
| RDS | Relational Database | Moderate to High | Varies by engine (MySQL, PostgreSQL, etc.), high read/write throughput with proper scaling | Use for traditional relational database workloads with complex queries and transactions | Supports multiple engines (MySQL, PostgreSQL, MariaDB, etc.), fully managed, scales vertically |
| AuroraDB | Relational Database | Moderate to High | High performance (5x MySQL, 3x PostgreSQL) | Use for high-performance relational workloads needing scalability and fault tolerance | Compatible with MySQL and PostgreSQL, auto-scaling, replicated storage across 3 AZs |
| DynamoDB | NoSQL (Key-Value Store) | Low to Moderate (based on read/write capacity) | Millisecond latency for reads and writes, scales horizontally | Use for low-latency applications with unstructured or semi-structured data | Serverless, fully managed, scales horizontally, great for key-value and document data models |
| DocumentDB | NoSQL (Document Store) | Moderate to High | High performance for JSON document workloads | Use for applications requiring MongoDB compatibility for document-based data | Managed MongoDB-compatible document store, scales horizontally with replica sets |
| Keyspaces | NoSQL (Wide-Column Store) | Low to Moderate (pay-per-request) | Low latency for high-throughput workloads | Use for Cassandra-compatible applications requiring massive scale and high availability | Managed Cassandra-compatible service, scales horizontally, great for wide-column data model |
| QLDB | Ledger Database | Moderate | High throughput, immutable and cryptographically verifiable transactions | Use for applications needing immutable, verifiable transaction logs (e.g., supply chain, finance) | Immutable ledger, ensures cryptographically verifiable transaction history |

| | | | | | |
|-------------|------------------------------------|--|---|--|---|
| Neptune | Graph Database | Moderate to High | Low-latency graph queries at scale, optimized for graph traversal | Use for graph-based applications like social networks, fraud detection, and recommendation engines | Supports both property graph (Gremlin) and RDF graph (SPARQL) models |
| Timestream | Time Series Database | Low to Moderate | Optimized for ingesting, storing, and querying time-series data | Use for IoT, DevOps, or industrial telemetry workloads requiring time-series data | Fully managed, scales automatically, optimized for time-series workloads |
| ElastiCache | In-Memory Cache (Redis, Memcached) | Low to High (depending on instance type) | Microsecond latency, great for caching | Use for caching, session management, and real-time analytics | Managed Redis or Memcached, great for improving application performance via caching |
| MemoryDB | In-Memory Database (Redis) | Moderate to High | Microsecond latency, optimized for high availability and durability | Use for real-time applications that need persistent in-memory data, such as gaming, chat, and leaderboards | Managed Redis with high availability and durability features |
| Redshift | Data Warehouse (SQL-based) | Moderate to High | Optimized for high-performance querying and reporting | Use for large-scale analytics, BI reporting, and complex SQL querying on petabyte-scale data | Columnar storage, massive parallel processing (MPP), great for large-scale data warehousing |

Cost

- RDS, AuroraDB, Redshift: Typically higher cost because of instance-based pricing and managed services with high performance.
- DynamoDB, Timestream: Pay-per-request pricing, making it cost-effective for low to medium traffic, but costs increase with higher throughput.
- ElastiCache, MemoryDB: Costs are based on the instance type and storage, with MemoryDB generally more expensive due to high availability and durability.
- Keyspaces, QLDB: Moderately priced, with Keyspaces offering pay-per-request pricing and QLDB pricing based on transaction volume and data storage.
- DocumentDB: Moderate to high cost due to the need for instance-based pricing and storage.

Performance Considerations

1. DynamoDB, ElastiCache, MemoryDB: Offer the highest performance due to their in-memory nature. High throughput with horizontal scaling.
2. AuroraDB, DocumentDB, Keyspaces, Neptune, Timestream: Provide high performance for most workloads.
3. RDS: Performance can vary depending on instance type and workload. Traditional relational database performance, but may require manual tuning.
4. QLDB: Provides high performance for ledger-based operations.
5. Redshift: Optimized for large-scale analytics.
6. Neptune: Optimized for graph queries, delivering low-latency graph traversal performance.

When to Use

- RDS/AuroraDB: For complex relational workloads with transactional support and complex joins/queries.
- DynamoDB: For highly scalable, low-latency key-value workloads.
- DocumentDB: For MongoDB-compatible applications that need a managed document database solution.
- QLDB: When you need cryptographic verification of transactional data.
- Neptune: For applications with complex relationships that require graph traversal and queries.
- Timestream: For IoT, time-series data, and telemetry that require fast ingestion and analysis of time-based data.
- ElastiCache/MemoryDB: For in-memory caching and real-time applications requiring microsecond-level access.

Additional Considerations:

- Aurora vs. RDS: Aurora offers better scalability and availability, especially for high-demand environments.
- DynamoDB vs. DocumentDB: DynamoDB is better for key-value workloads, whereas DocumentDB is more suitable for document-based, MongoDB-compatible use cases.
- Timestream vs. QLDB: Timestream is best for time-series data, whereas QLDB is for immutable, verifiable transaction logs.
- ElastiCache vs. MemoryDB: MemoryDB is for real-time, durable, in-memory data use cases, while ElastiCache is more for caching and session management.